



Heap Exploitation

Binary Exploitation 2019/2020

Heap Overview



- The Heap is a pool of memory used for dynamic allocation at runtime.
 - `malloc` – grabs memory from the heap
 - `free` – releases memory on the heap
- It is a data segment, just like the `.stack` or the `.bss`

Heap Overview

```
pwndbg> vmmap
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
0x555555554000 0x555555555000 r-xp 1000 0 /home/vagrant/share/heap/hi_mom
0x5555555754000 0x5555555755000 r--p 1000 0 /home/vagrant/share/heap/hi_mom
0x5555555755000 0x5555555756000 rw-p 1000 1000 /home/vagrant/share/heap/hi_mom
0x5555555756000 0x5555555777000 rw-p 21000 0 [heap]
0x7ffff79e4000 0x7ffff7bcb000 r-xp 1e7000 0 /lib/x86_64-linux-gnu/libc-2.27.so
0x7ffff7bcb000 0x7ffff7dcb000 ---p 200000 1e7000 /lib/x86_64-linux-gnu/libc-2.27.so
0x7ffff7dcb000 0x7ffff7dcf000 r--p 4000 1e7000 /lib/x86_64-linux-gnu/libc-2.27.so
0x7ffff7dcf000 0x7ffff7dd1000 rw-p 2000 1eb000 /lib/x86_64-linux-gnu/libc-2.27.so
0x7ffff7dd1000 0x7ffff7dd5000 rw-p 4000 0
0x7ffff7dd5000 0x7ffff7dfc000 r-xp 27000 0 /lib/x86_64-linux-gnu/ld-2.27.so
0x7ffff7dd5000 0x7ffff7dfc000 rwxp 27000 0 <explored>
0x7ffff7fe9000 0x7ffff7feb000 rw-p 2000 0
0x7ffff7ff7000 0x7ffff7ffa000 r--p 3000 0 [vvar]
0x7ffff7ffa000 0x7ffff7ffc000 r-xp 2000 0 [vdso]
0x7ffff7ffc000 0x7ffff7ffd000 r--p 1000 27000 /lib/x86_64-linux-gnu/ld-2.27.so
0x7ffff7ffd000 0x7ffff7ffe000 rw-p 1000 28000 /lib/x86_64-linux-gnu/ld-2.27.so
0x7ffff7ffe000 0x7ffff7fff000 rw-p 1000 0
0x7ffffffffffde000 0x7ffffffffff000 rw-p 21000 0 [stack]
0xffffffffffff600000 0xffffffffffff601000 r-xp 1000 0 [vsyscall]
```

Heap Overview - Usage



```
int main() {  
    char * buffer = NULL;  
  
    /* get 0x100 bytes from memory from the heap */  
    buffer = (char*) malloc(sizeof(char) * 0x100);  
  
    fgets(stdin, buffer, 0x100);  
    printf("%s", buffer);  
  
    /* release the allocated memory */  
    free(buffer);  
  
    return 0;  
}
```

Heap Chunks



- The heap is made of **Heap Chunks**, and there are different types:
 - Allocated chunk
 - Free chunk
 - Top chunk
 - Last Remainder chunk

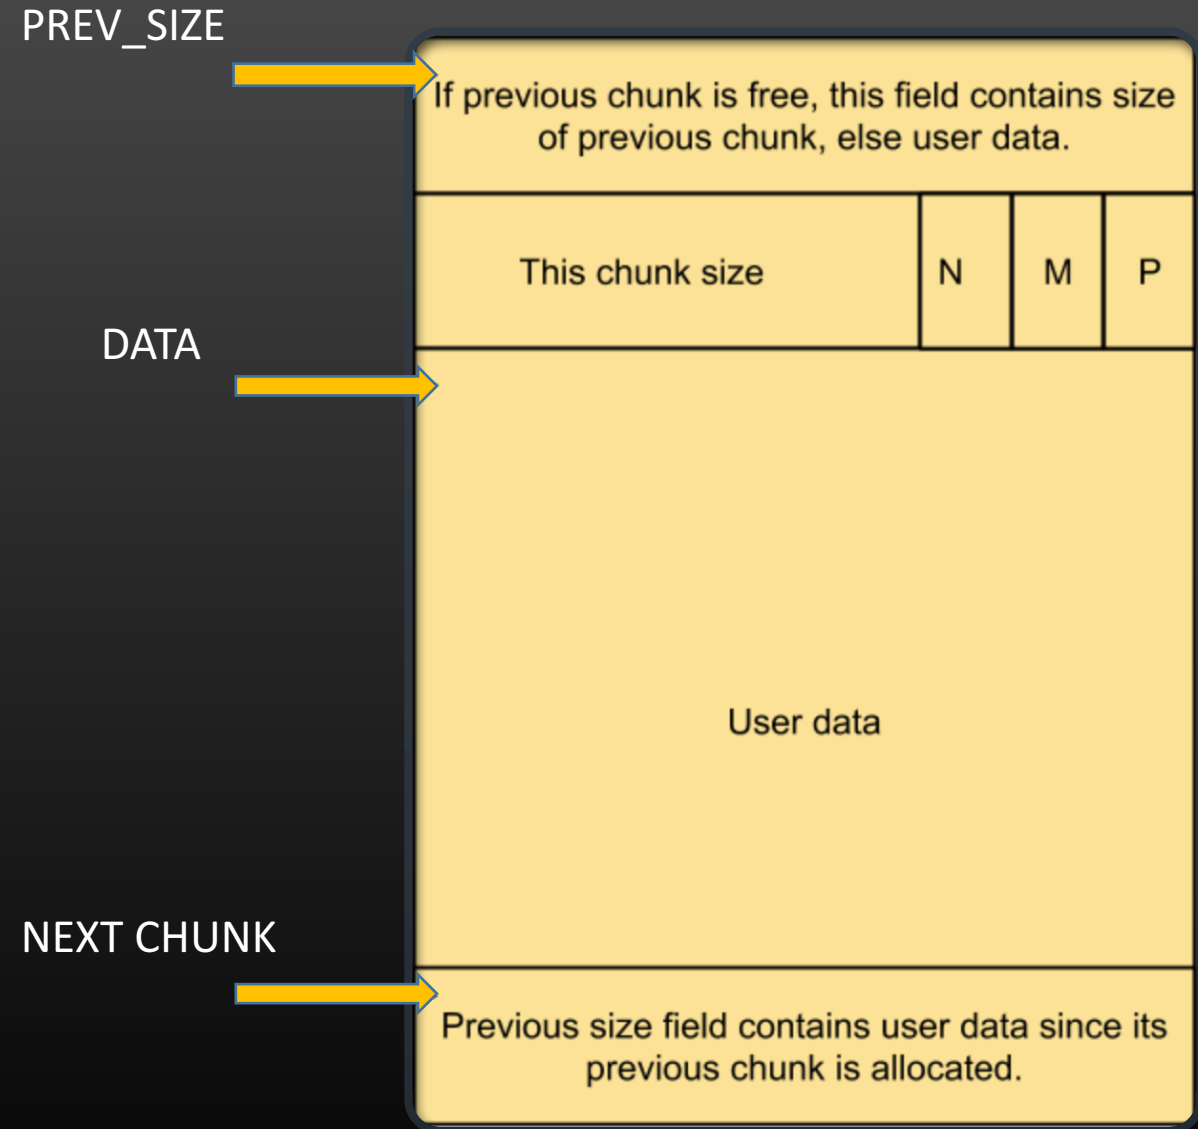


Chunk Implementation

```
struct malloc_chunk {  
    INTERNAL_SIZE_T  prev_size;  /* Size of previous chunk (if free). */  
    INTERNAL_SIZE_T  size;       /* Size in bytes, including overhead. */  
  
    struct malloc_chunk* fd;      /* double links -- used only if free. */  
    struct malloc_chunk* bk;  
  
    /* Only used for large blocks: pointer to next larger size. */  
    struct malloc_chunk* fd_nextsize;  
    struct malloc_chunk* bk_nextsize; /* double links -- used only if free. */  
}
```

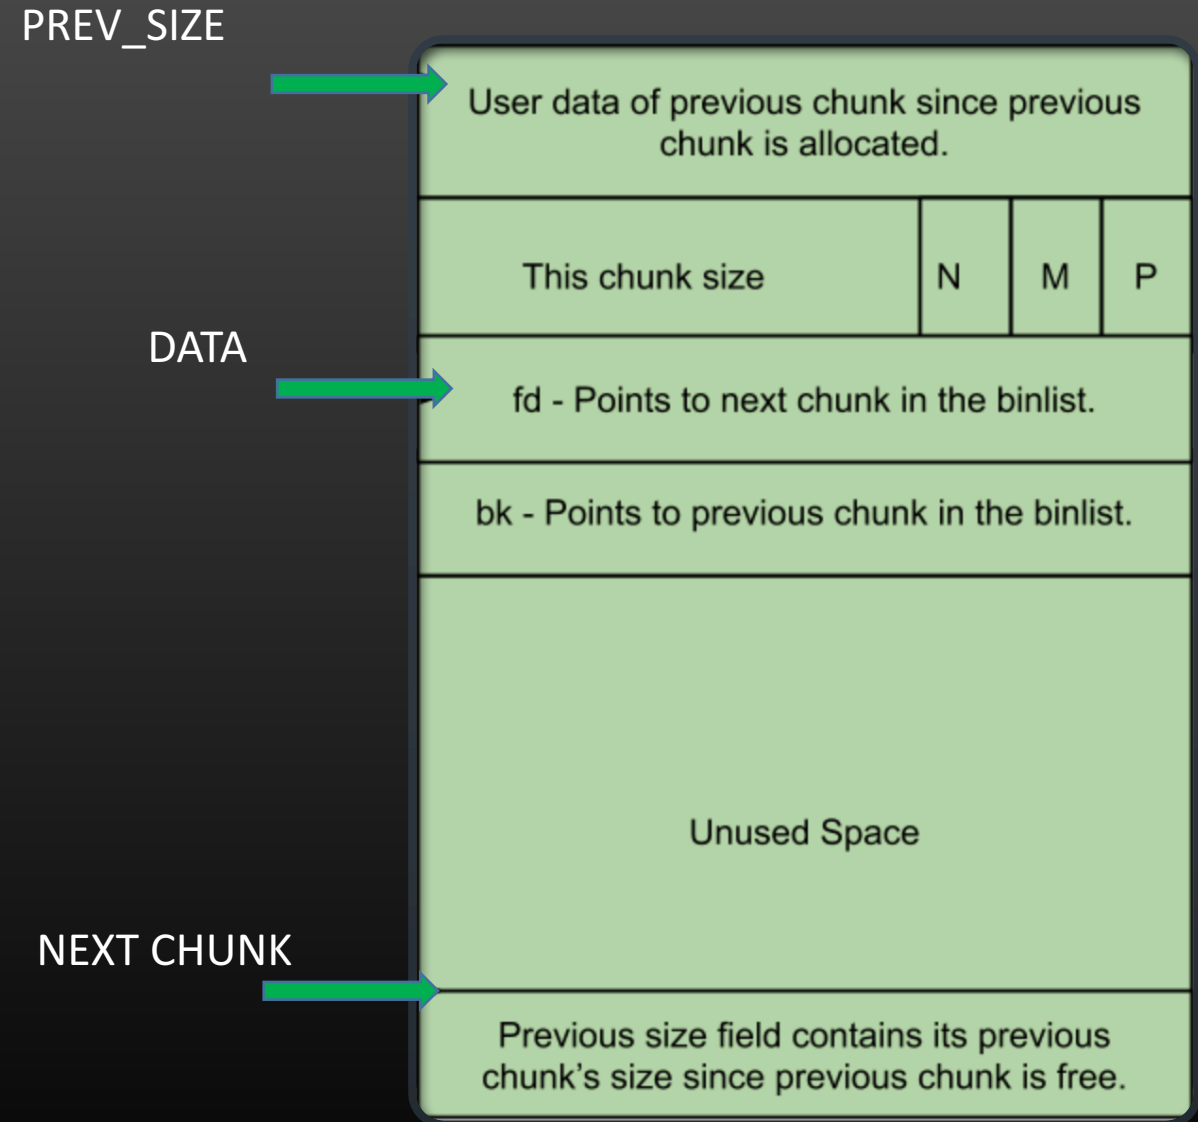
Allocated Chunk

- **PREV_SIZE**: size of previous chunk if allocated.
- **SIZE**: size of the chunk. The 3 lsbits are ignored (always 0) for size purposes.
- **PREV_INUSE(P)**: bit is set if previous chunk is in use
- **IS_MMAPPED(M)** – bit set if chunk was mmap'd
- **NON_MAIN_ARENA (N)** – bit set when chunk belongs to a thread arena



Free Chunk

- **PREV_SIZE**: previous chunk's user data
- **SIZE**: size of the chunk. The 3 lsbits are ignored (always 0) for size purposes.
- **PREV_INUSE(P)**: bit is set if previous chunk is in use
- **IS_MMAPPED(M)** – bit set if chunk was mmap'd
- **NON_MAIN_ARENA (N)** – bit set when chunk belongs to a thread arena





Top Chunk

- Used to service user request when there are **NO FREE CHUNKS** in any bin.
- Features:
 - If `top_chunk->size > requested->size`, it is split in two
 - User chunk (requested size)
 - Remainder chunk (of remaining size)
 - If not, the top chunk is extended using the `sbrk` or `mmap` syscalls.



Last Remainder Chunk

- This chunk is the remainder from the most recent split
- Why?
 - Helps to improve locality of reference, increasing performance

Coalescing



- Two chunks which are free can't be adjacent, so they are combined into a single **Free chunk**.
- Why?
 - It eliminates external fragmentation, but it slows up free!!!



Main Arena - Bins

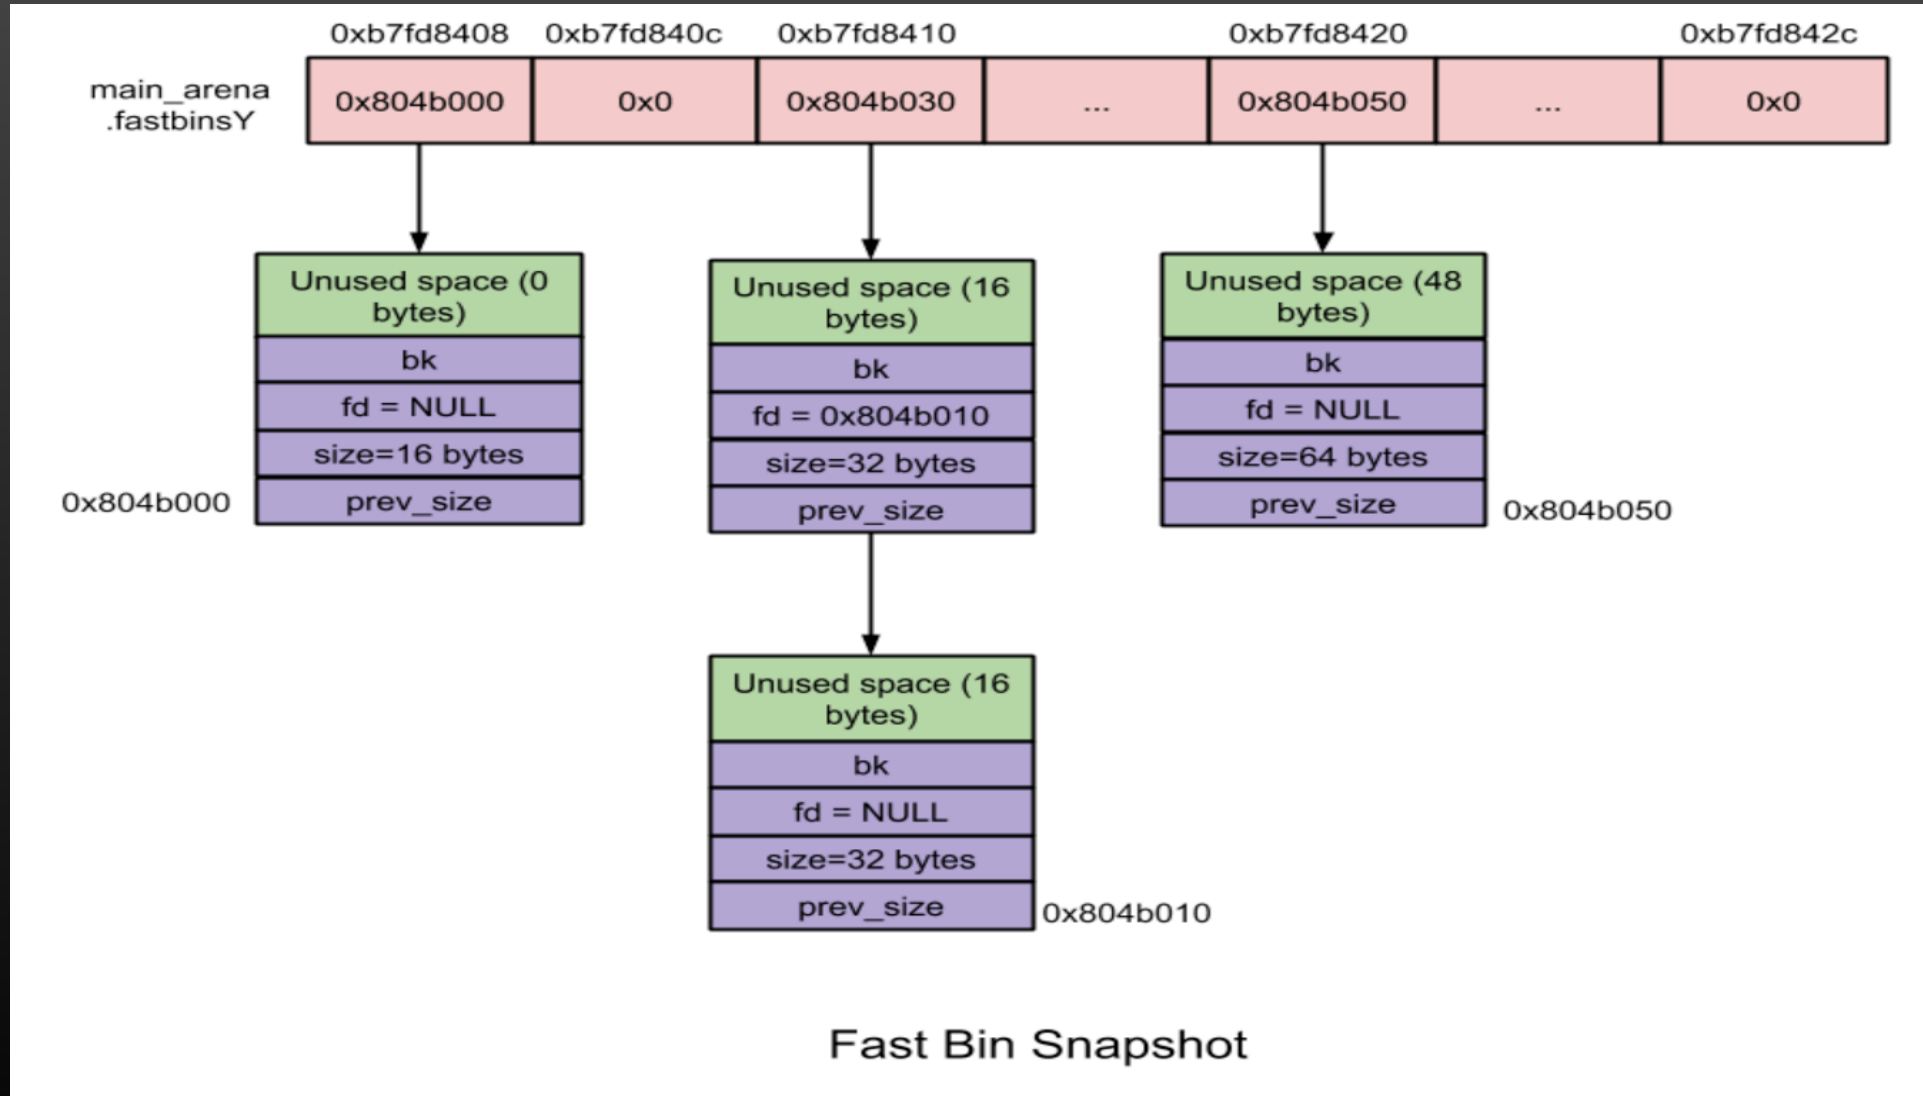
- Bins are the free-list data structures.
- They are used to hold free chunks.
- Based on the chunk size, different bins are available:
 - Fast bin
 - Unsorted bin
 - Small bin
 - Large bin

Fast Bins



- Singly-linked list
- 16 bytes \leq `chunk->size` \leq 80 bytes
- Non Coalescing
 - Chunks that belong to a Fast bin don't coalesce

Fast Bins



Small Bins



- Doubly-linked list
- 80 bytes < `chunk->size` <= 512 bytes
- Coalescing

Large Bins



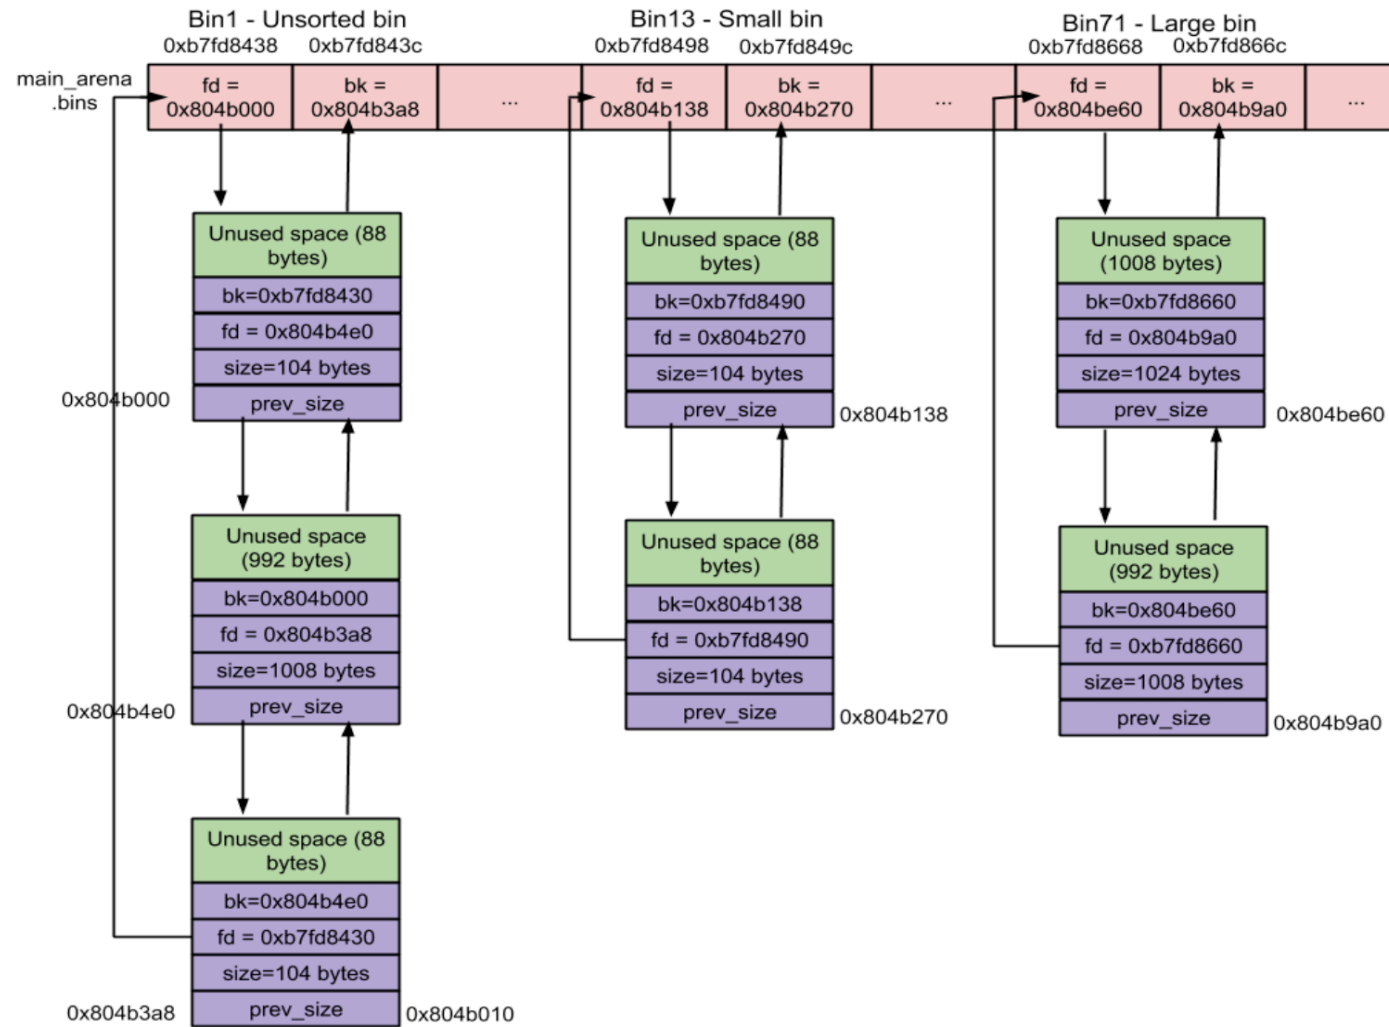
- Doubly-linked list
- 512 bytes < chunk->size
- Coalescing

Unsorted Bin



- When small or large chunks get freed, they are added to the Unsorted Bin
- Why?
 - Helps speed up memory allocation
- Doubly-linked list
- No size requirement

Main Arena - Bins



Unsorted, Small and Large Bin Snapshot

Tcache



- Structure introduced in glibc 2.26 in order to improve the heap's performance.
- The Tcache is a per-thread structured, stored on the heap, and it's very similar to the fast bin's structure.



Tcache Structure

```
typedef struct tcache_entry {  
    struct tcache_entry *next;  
} tcache_entry;
```

```
typedef struct tcache_perthread_struct {  
    char counts[TCACHE_MAX_BINS];  
    tcache_entry *entries[TCACHE_MAX_BINS];  
} tcache_perthread_struct;
```

Put a chunk in tcache



static void

```
tcache_put (mchunkptr chunk, size_t tc_idx) {  
    tcache_entry *e = (tcache_entry *) chunk2mem (chunk);  
    assert (tc_idx < TCACHE_MAX_BINS);  
    e->next = tcache->entries[tc_idx];  
    tcache->entries[tc_idx] = e;  
    ++(tcache->counts[tc_idx]);  
}
```

Get chunks from tcache



```
static void *  
tcache_get (size_t tc_idx) {  
    tcache_entry *e = tcache->entries[tc_idx];  
    assert (tc_idx < TCACHE_MAX_BINS);  
    assert (tcache->entries[tc_idx] > 0);  
    tcache->entries[tc_idx] = e->next;  
    --(tcache->counts[tc_idx]);  
    return (void *) e;  
}
```



Historical Vulnerabilities and Attacks

- Use After Free
- Double Free
- Poison null byte
- Unsafe Unlink
- Fastbin dup/poison
- Tcache dup/poison